

# Training on Log-based Field Failure Data Analysis

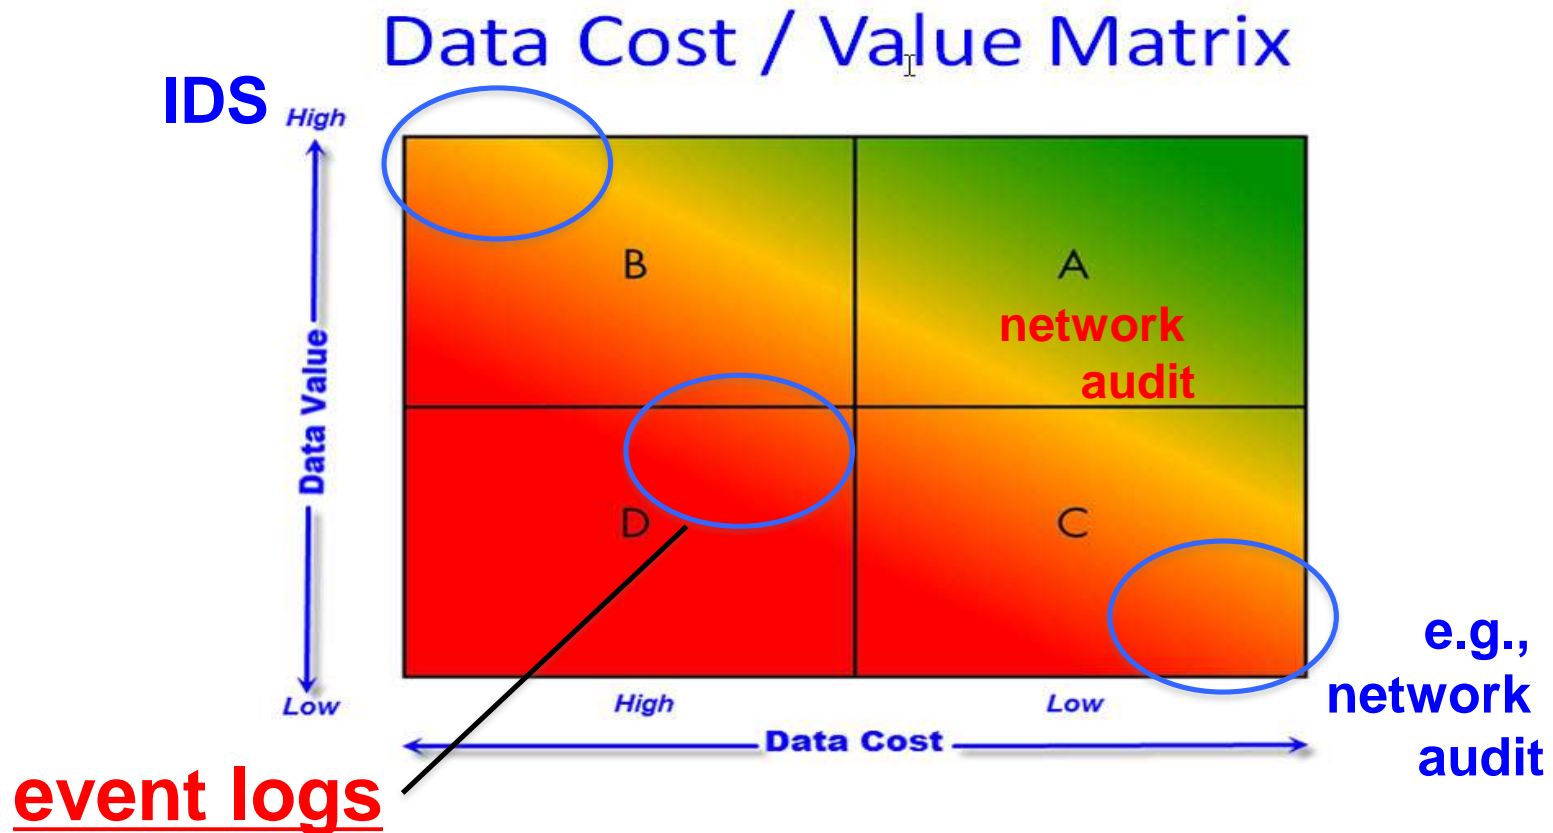
Domenico Cotroneo  
cotroneo@unina.it

# Data-driven Analysis

- Textual logs and numeric data produced by applications, systems, and IDSs are a key resource to:
  - **characterize failures;**
  - predict problems;
  - pinpoint security attacks;
  - improve the system;
- Data are available in many critical systems!




# Not all the data is the same



# Challenges

- **Dealing with operational data is extremely hard:**
  - increasing network traffic and failure rates;
  - increasing data volumes;
  - data quality and heterogeneity;
  - lack of systematic procedures;
  - low automation degree;
  - ...

**5min in-out traffic  
of a data center**



# What is the training lesson about

- Conducting a **log-based FFDA** of complex, large-scale systems.
- Real data collected from:
  - **Mercury** at the National Center for Supercomputing Application (NCSA at University of Illinois);
  - **BlueGene/L** at the Lawrence Livermore National Labs (LLNL, CA, USA).

# Goals

- **Objectives:** to learn the basics in
  - data classification/analysis;
  - determining the coalescence window to group a set of error entries;
  - performing a failure analysis;
  - reliability modelling via well-know distributions;
  - analysis at system and node level;
  - ...

# Why is it hard?

- Monitoring in a **datacenter day** ...





- Even if the data is available, the real challenge is still being actually able to use it!





# Why?


- Failure characterization;
- Failure correlation/localization;
- Failure prediction;
- Cause analysis;
- Check-pointing tuning;
- Failure-aware job allocation;
- ...

**Correct understanding  
of the failure  
behavior is a critical task!**



# Mercury cluster: overview

- Mercury consists of IBM nodes. The cluster has a 3-layer architecture + a management node (**tg-master**)



- Myrinet. The nodes run a RedHat 9.0 OS
- Entries have been collected via the **syslog daemon**

# Mercury event log

- 80,854 log **filtered and anonimized** entries (*only FATAL error entries are given in the log*)
- Format:
  - Timestamp
  - Originating node
  - Originating subsystem
  - Text-free message

```
1167667229 tg-c238 DEV
```

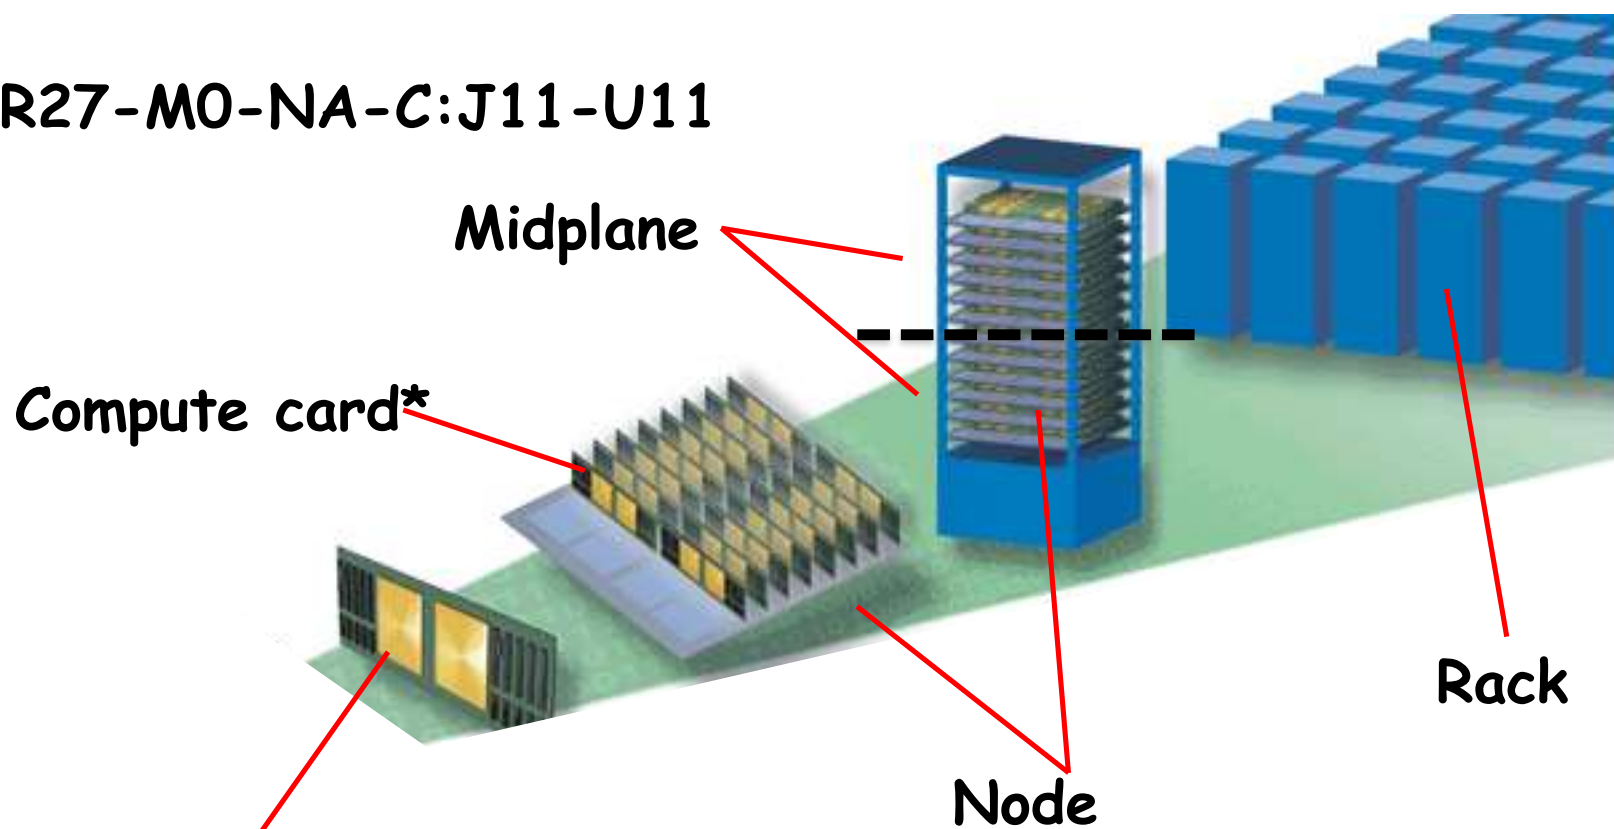
```
Component Info: Vendor Id =* *, Device Id =* *, Class Code =*  
*, Seg Bus Dev Func =* * * * *
```

# Examples of error entries

- [DEV]  
+Platform PCI Component Error Info Section
- [MEM]  
Mem Error Detail: Physical Address \*, Address Mask:\*, Node:\*,  
Card:\*, Module: \*, Bank:\*, Device:\*,Row:\*, Column:\*
- [NET]  
connection down
- [I-O] - error, dev \*:\*(hda), sector \*  
- hda: packet command error: error=\*
- [PRO]  
+BEGIN HARDWARE ERROR STATE AT CMC

# BG/L Sample Diagram

R27-M0-NA-C:J11-U11



Compute chip

\*16 cards / node (J02-J17)  
for each midplane **N0, N4, N8, NC**  
**have an additional I/O card (J18)**

# BG/L event log\*

- 125,624 **filtered** log entries (*only FATAL error entries are given in the log*)
- Format:
  - Timestamp
  - Originating node
  - **Originating card**
  - Text-free message

1128621351 R04-M0-N0 **J18-U01**

Lustre mount FAILED : bgllo66 : block\_id : location

---

\*Log made publicly available by J. Stearley from Sandia National Labs  
<http://www.cs.sandia.gov/~jrstear/logs/bgl.gz>



# Examples of error entries

1135619785 R72-M1-NC J09-U11 parity error in read queue 0.....0

1135780093 R35-M0-N3 J12-U01 machine check interrupt (bit=0x1d): L2  
dcache unit write data parity error

1132173448 R46-M1-N8 J04-U11 Error receiving packet on tree network,  
expecting type 57 instead of type 3 (softhead=0030114e 602d0003  
00000002 00000000) PSR0=10011f01 PSR1=00000000 PRXF=00000002  
PIXF=00000007

...

# Tools usage and sample analysis\*

---

\*Results shown in the rest of the presentation are based on a **subset** of 10,000 entries of Mercury log (MercuryErrorLogTEST.txt)

# Tools needed for the analysis


- A suite of home-made **BASH scripts** (tuple count, grouping of the error entries; total 4 scripts)
  - the basic scripts can be used sequentially to perform the FFDA of the log;
  - the usage of the scripts is described in the presentation;
  - **students are free to modify/improve or to write further scripts to automate the analysis;**
  - a baseline **ffdatoolset.zip** archive is provided to the students.
- Matlab & **statistics toolbox** (TTF estimation, curve fitting, ...)

# How many failures?

- The same problem is reported multiple times in the log
- Entries that are possibly related to the same problem **must be grouped** (or coalesced).

```
$ head -n 10 MercuryErrorLogTEST.txt
1167637660 tg-c645 PRO  tg-c645__+BEGIN HARDWARE ERROR STATE AT CMC
1167637660 tg-c645 PRO  tg-c645__Device Error Info Section
[... omissis ...]
1167637720 tg-c645 PRO  tg-c645__Device Error Info Section
1167637720 tg-c645 PRO  tg-c645__Error Map: *
1167655228 tg-c238 DEV  tg-c238__+BEGIN HARDWARE ERROR STATE AT CPE
1167655228 tg-c238 DEV  tg-c238__+Platform PCI Component Error Info Section
[... omissis ...]
```

# How many failures?



- The same fault can be activated multiple times and generate multiple errors that propagate within the system.

# Determining the coalescence window



# How many failures?

- The script `tupleCount_func_CWIN.sh` analyzes how the tuple count varies for increasing values of the *coalescence window*
- Execution takes several min/hour depending on the number of entries in the log\*
- The output (CWIN,tuple-count pairs) is saved in the `tupleCount-<log_filename>.txt` file

```
$/tupleCount_func_CWIN.sh MercuryErrorLogTEST.txt tentative-Cwin.txt
```

```
$ cat tupleCount-logMERCURY-TEST.txt
```

```
1      2049
```

```
2      1238
```

```
3      1063
```

```
4      1033
```

```
[... omissis ...]
```

CWIN size

count

\*around 1hour for the `MercuryErrorLogTEST.txt` log; estimated around 6/7hours for the *exercise* log (run with an Intel Core2 Duo @2.4Ghz processor)


# Determining the "optimal" CWIN

- Plotting the tuple-count curve

```
>> load tupleCount-MercuryErrorLogTEST.txt  
>> plot(tupleCount_MercuryErrorLogTEST(:,1),tupleCount_MercuryErrorLogTEST(:,2),'-*b');  
>> xlabel('CWIN [s]');ylabel('count');  
>> legend('tuple count');
```

- Rule: a *good CWIN* is right after the "knee" of the curve

CWIN: 240sec (4min)  
276 tuples




# Grouping the entries with the selected CWIN

- The script `tupling_with_CWIN.sh` coalesces the entries with the selected CWIN (the script takes around 2min).
- Output (summary files + content of each individual tuple) is saved in the directory `tupling_<log-filename>-<CWIN>`

```
$ ./tupling_with_Cwin.sh MercuryErrorLogTEST.txt 240
*** Grouping the entries for CWIN=240 ***
  - creating tuple # 1
  - creating tuple # 2
  [...]
$ cd tupling_MercuryErrorLogTEST-240/
$ ls
interarrivals.txt  tuple_16      tuple_30      tuple_45
lengths.txt       tuple_17      tuple_31      tuple_46
startingPoints.txt tuple_18      tuple_32      tuple_47
tuple_1           tuple_76      tuple_90      [...]
```

# Output files



- $L_1, L_2, \dots, L_N$  tuple lengths (*lengths.txt*)
- $I_1, I_2, \dots, I_{N-1}$  interarrivals (*interarrivals.txt*)
- $S_1, S_2, \dots, S_N$  timestamps of the init point of the tuples (*1st column of startingPoints.txt*)

\* NOTE: each interarrival is a TTF sample \*

# Issue #1: truncations

- A single failure might be broken into multiple tuples

## Tuple #4

...

1167657550 tg-c238 MEM Physical Address x, Address Mask: x, Node: x, Card: x, Module: x, Bank: x, Device: x, Row: x, Column: x,

1167657550 tg-c238 MEM Physical Address x, Address Mask: x, Node: x, Card: x, Module: x, Bank: x, Device: x, Row: x, Column: x,


## Tuple #5

1167657941 tg-c238 DEV +BEGIN HARDWARE ERROR STATE AT CPE

1167657941 tg-c238 DEV +Platform PCI Component Error Info Section


- Time distance between tuples is 390sec

# Issue #2: collisions






# Issue #2: collisions



# Issue #2: collisions

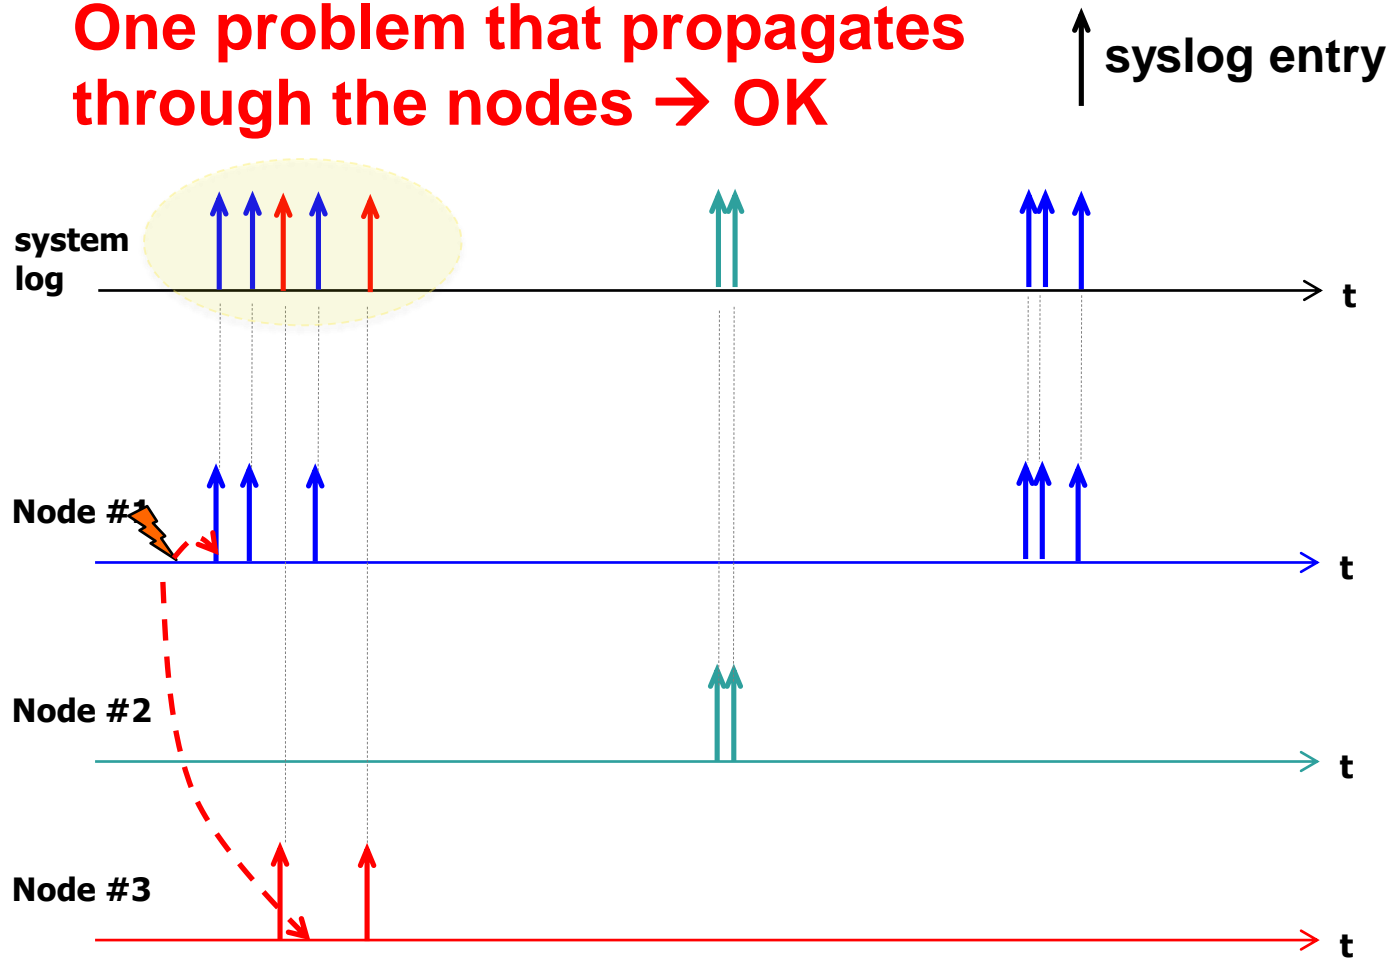
Two independent, overlapping,  
Problems → COLLISION

↑ syslog entry



# Issue #2: collisions

One problem that propagates through the nodes → OK



# Example

- Tuple #61 includes entries produced by 7 nodes! The tuple lasts 1244 sec (20 min)
- Are the events correlated?

3	tg-c669
3	tg-login1
144	tg-login2
144	tg-login3
2	tg-login4
8	tg-s131
28	tg-s159

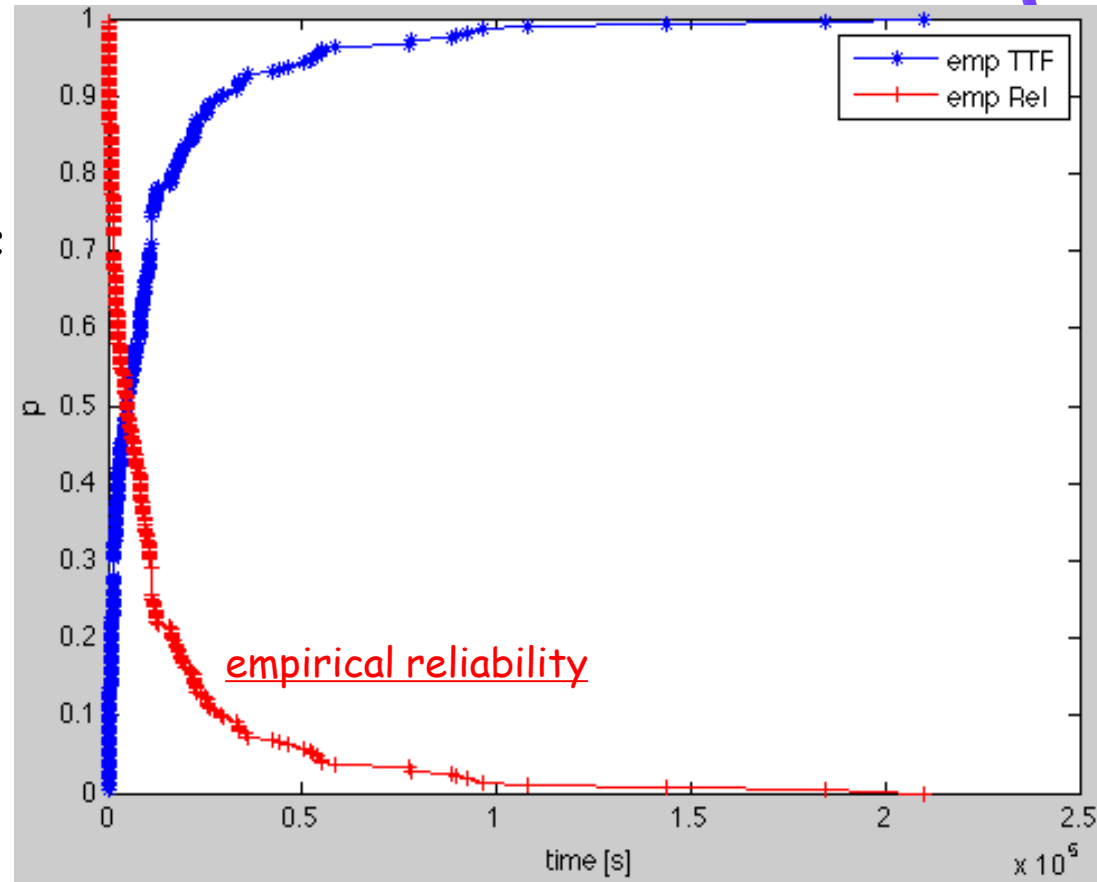
# Reliability modeling

# TTF: empirical distribution

- The **interarrivals.txt** file is imported in Matlab to obtain the *empirical CDF* of the TTF of the system.

```
>> load interarrivals.txt
>> [y,t]=cdfcalc(interarrivals);
>> empTTF=y(2:size(y,1));
>> empRel=1-empTTF;
>> plot(t,empTTF,'-*b',t,empRel,'-+r');
>> xlabel('time [s]'); ylabel('p');
>> legend('emp TTF', 'emp Rel');
```

- Given  $t$ ,  $TTF(t)$  is the probability that the system exhibits a failure after  $t$ sec since the last fail. occurred.





# Fitting the empirical reliability

- The cftool (statistics toolbox) can be used to fit a theoretical distribution (thRel) to the empirical (empRel) one (“Fitting” button)


```
>> cftool(t,empRel);
```

- Tentative type of distribution:

exponential

$$\text{thRel}(t) = ae^{bt} \quad a \geq 0$$

$$b < 0$$



# Using the "Fit Editor"

- 1) Press "New fit"
- 2) Give a name to the fit
- 3) Select the type of fit (e.g., exponential)
- 4) Select the equation
- 5) Press "Apply"
- 6) Save the results to Matlab workspace

The Fit Editor dialog box is shown with the following configuration:

- Fit name: thRel1
- Data set: empRel vs. t
- Exclusion rule: (none)
- Type of fit: Exponential
- Equation:  $a \cdot \exp(b \cdot x)$
- Center and scale X data: ☐

Buttons: New fit, Copy fit, Fit options..., Immediate apply, Cancel, Apply.

**Results**

General model Exp1:  
 $f(x) = a \cdot \exp(b \cdot x)$   
Coefficients (with 95% confidence bounds):  
a = 0.8898 (0.8773, 0.9023)  
b = -0.0001054 (-0.0001096, -0.0001012)

Goodness of fit:  
SSE: 0.6735  
R-square: 0.9689  
Adjusted R-square: 0.9688  
RMSE: 0.05159

**Table of Fits**

Fit name	Data set	Equation name	SSE	R-square
thRel1	empRel vs. t	Exp1	0.67345364104988...	0.96891204628...

Buttons: Delete fit, Save to workspace..., Table options..., Close, Help.

Coefficients (a and b)

Quality of the fit  
(SSE, R-square)

Save Fit to MATLAB Workspace

☒ Save fit to MATLAB object named: thRel1

☐ Save goodness of fit to MATLAB struct named:

☐ Save fit output to MATLAB struct named:

Buttons: OK, Cancel.

# Assessing the quality of the fit

- *Is the theoretical distribution a good fit for the empirical one?*


- Indicators:

**SSE**: Sum of Squares due to Errors (has to be close to 0)

**R-square**: “ability” of the fit to explain the variation in the data

thRel1:                      SSE=0.6735  
                                    R-square=0.9688

- **Kolmogorov-Smirnov test**



*Analysis at node/  
error-category level*

# Overview of the log

- The script `logStatistics.sh` provides an overview of the data in the log (useful to figure out entries-prone categories and/or nodes)
- distribution of the entries by **error category**
- **nodes** with the largest number of entries

```
$ ./logStatistics.sh MercuryErrorLogTEST.txt
```

```
== Total error entries ==  
10000
```

```
== Breakup by CATEGORY ==
```

```
DEV 5576  
PRO 1292  
MEM 1198  
I-O 1195  
NET 713  
OTH 26
```

```
== Breakup by NODE* ==
```

```
tg-c572 4030  
tg-c238 1272  
tg-master 817  
tg-c242 622  
tg-c648 616  
tg-c550 253  
tg-c894 217  
tg-c284 180  
tg-c451 159  
tg-s176 156
```

```
* only the 10 most occurring nodes are reported
```

# Filtering the log

- The script **filter.sh** is then used to extract entries (by node or category) from the log.

Output is saved in a file named **<filterName>-<systemLog>.txt**

logFile

filterName

```
$/filter.sh MercuryErrorLogTEST.txt tg-master  
Filtering by NODE: tg-master ... [DONE]  
$/filter.sh MercuryErrorLogTEST.txt DEV  
Filtering by CATEGORY/CARD: DEV ... [DONE]
```

- Filtering can be done either by node or by category*
- Filtering can be **sequential**: e.g., all MEM errors of tg-c572  
\$/filter.sh MercuryErrorLog.txt tg-c572  
\$/filter.sh tg-c572-MercuryErrorLog.txt MEM  
**result is MEM-tg-c572-MercuryErrorLog.txt**

# Error distribution by node

- `logStatistics.sh` applied to a single node returns the most-occurring categories for that node

```
./logStatistics.sh tg-master-MercuryErrorLogTEST.txt
```


```
== Total error entries ==  
817
```

```
== Breakup by CATEGORY ==  
NET 663  
I-O 152  
OTH 1  
DEV 1
```

```
== Breakup by NODE* ==  
tg-master 817  
* only the 10 most occurring nodes are reported
```

# Example: error categories


```
./filter.sh MercuryErrorLogTEST.txt PRO
./filter.sh MercuryErrorLogTEST.txt I-O
./tupleCount_func_CWIN.sh PRO-MercuryErrorLogTEST.txt tentative-Cwin.txt
./tupleCount_func_CWIN.sh I-O-MercuryErrorLogTEST.txt tentative-Cwin.txt
```





# Example: error categories

```
$/tupling_with_Cwin.sh I-O-MercuryErrorLogTEST.txt 360  
$/tupling_with_Cwin.sh PRO-MercuryErrorLogTEST.txt 90
```



# Exercises

- Repeat the analysis for **MercuryErrorLog.txt** and **BGLErrorLog.txt**:
  - tuple count plot + selected CWIN;
  - grouping the entries with CWIN;
  - empirical distributions;
  - fit the empirical reliability (exponential model + other tentative distributions, e.g., **weibull**);
  - analysis of the fit via the K-S test;

**COMPARE RESULTS ACROSS THE SYSTEMS!**

# Exercises

- Conduct the analysis to answer the following questions:
  - can the same coalesce windows be used across different nodes (Mercury, BG/L) and error categories (Mercury)?
  - what about the relationship between the system reliability and nodes reliability?
  - there exist dependability-bottlenecks (i.e., top-contributors to the total number of failures)?
  - do similar functional nodes (e.g., two Mercury tg-cX computing nodes or BG/L I/O nodes Ri:Mx:Nz) exhibit similar reliability parameters?
  - there exist a relationship between the error types and the node (Mercury)?
  - Whatever additional considerations and results!